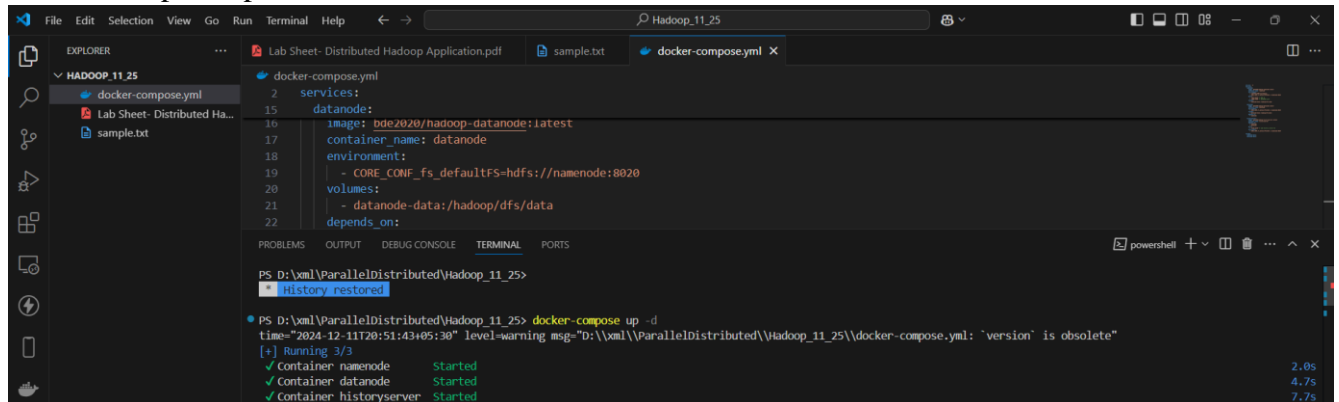


SE6103 Parallel and Distributed Systems

Lab Sheet: Distributed Hadoop Application

19APSE4325 – Tharindu Lakmal

1. docker-compose.yml
2. docker-compose up -d

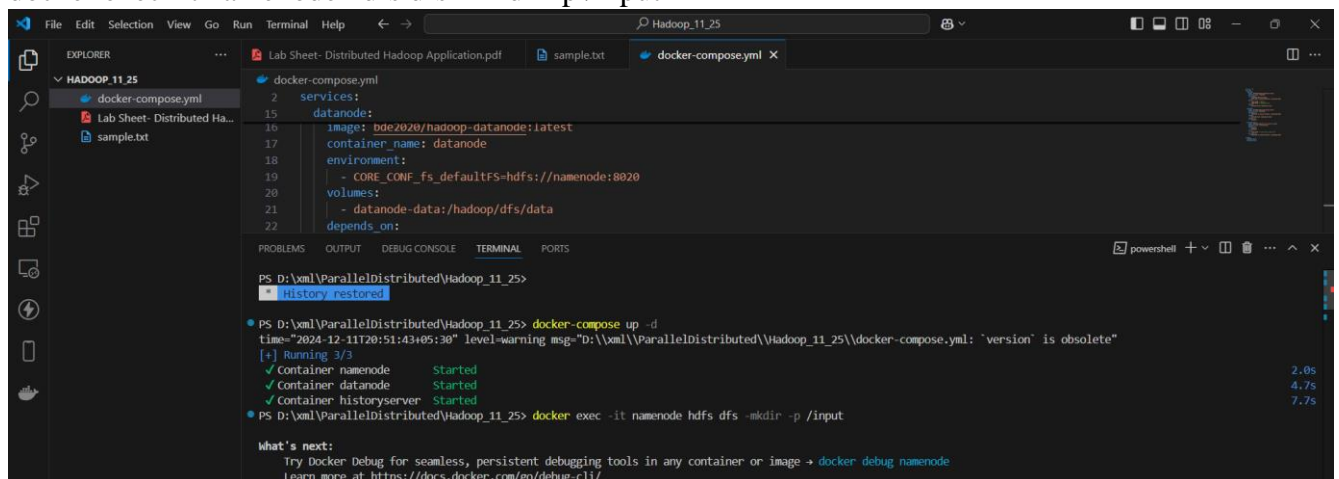


```
EXPLORER
  HADOOP_11_25
    docker-compose.yml
    Lab Sheet- Distributed Ha...
    sample.txt

  docker-compose.yml
  2  services:
  15  datanode:
  16    image: bde2020/hadoop-datanode:latest
  17    container_name: datanode
  18    environment:
  19      - CORE_CONF_fs_defaultsfs-hdfs://namenode:8020
  20    volumes:
  21      - datanode-data:/hadoop/dfs/data
  22    depends_on:

  PROBLEMS  OUTPUT  DEBUG CONSOLE  TERMINAL  PORTS
  PS D:\vml\ParallelDistributed\Hadoop_11_25>
  History restored
  PS D:\vml\ParallelDistributed\Hadoop_11_25> docker-compose up -d
  time="2024-12-11T20:51:43+05:30" level=warning msg="D:\vml\ParallelDistributed\Hadoop_11_25\docker-compose.yml: 'version' is obsolete"
  [+] Running 3/3
  ✓ Container namenode Started 2.0s
  ✓ Container datanode Started 4.7s
  ✓ Container historyserver Started 7.7s
```

3. docker exec -it namenode hdfs dfs -mkdir -p /input



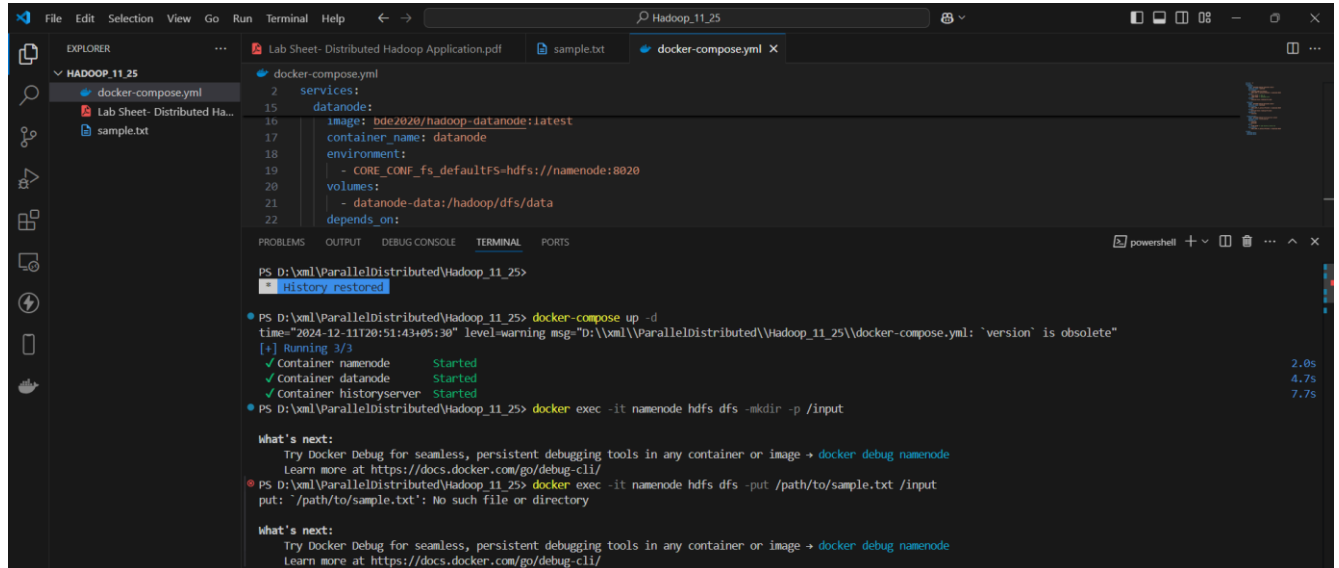
```
EXPLORER
  HADOOP_11_25
    docker-compose.yml
    Lab Sheet- Distributed Ha...
    sample.txt

  docker-compose.yml
  2  services:
  15  datanode:
  16    image: bde2020/hadoop-datanode:latest
  17    container_name: datanode
  18    environment:
  19      - CORE_CONF_fs_defaultsfs-hdfs://namenode:8020
  20    volumes:
  21      - datanode-data:/hadoop/dfs/data
  22    depends_on:

  PROBLEMS  OUTPUT  DEBUG CONSOLE  TERMINAL  PORTS
  PS D:\vml\ParallelDistributed\Hadoop_11_25>
  History restored
  PS D:\vml\ParallelDistributed\Hadoop_11_25> docker-compose up -d
  time="2024-12-11T20:51:43+05:30" level=warning msg="D:\vml\ParallelDistributed\Hadoop_11_25\docker-compose.yml: 'version' is obsolete"
  [+] Running 3/3
  ✓ Container namenode Started 2.0s
  ✓ Container datanode Started 4.7s
  ✓ Container historyserver Started 7.7s
  PS D:\vml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -mkdir -p /input

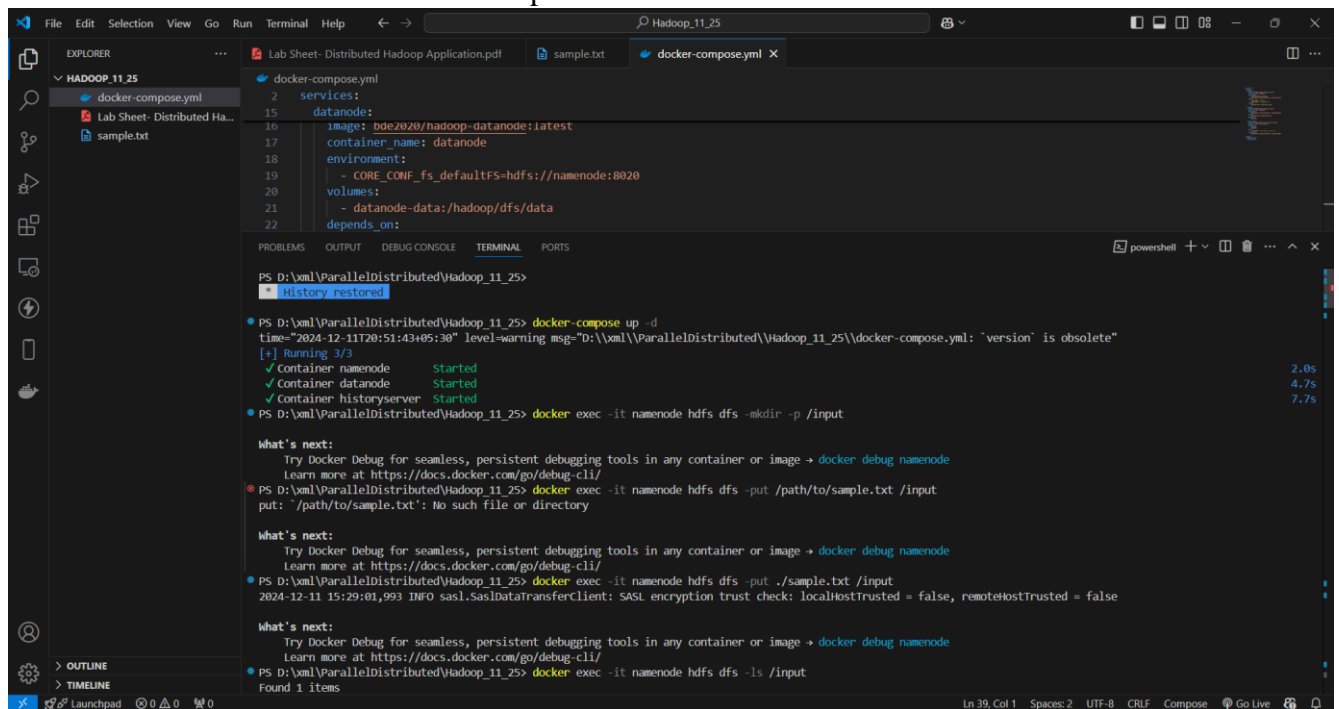
  What's next:
  Try Docker Debug for seamless, persistent debugging tools in any container or image → docker debug namenode
  Learn more at https://docs.docker.com/go/debug-cli/
```

4. `docker exec -it namenode hdfs dfs -put /path/to/sample.txt /input`

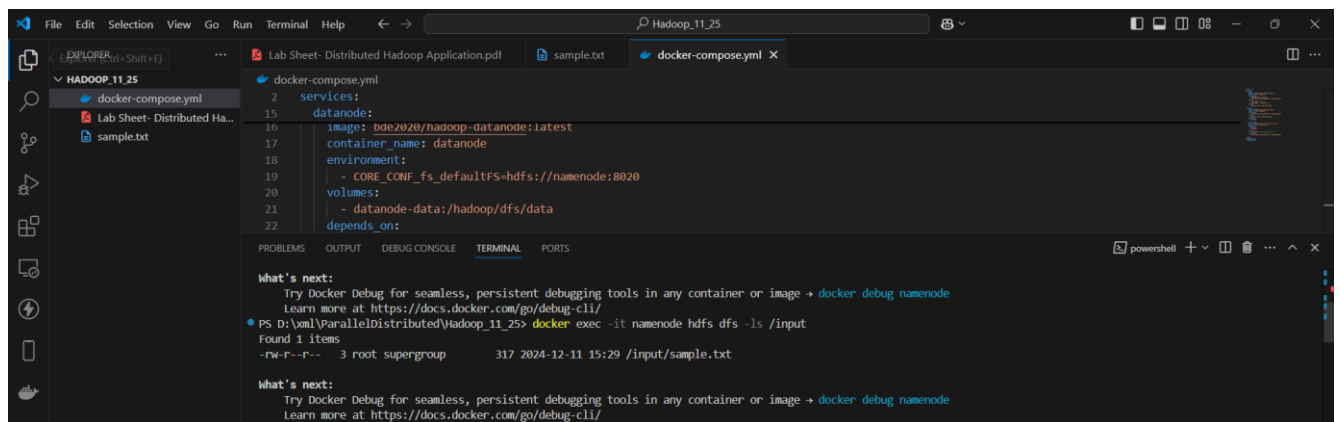


```
File Edit Selection View Go Run Terminal Help
Hadoop_11_25
EXPLORER
  HADOOP_11_25
    docker-compose.yml
    Lab Sheet- Distributed Ha...
    sample.txt
  docker-compose.yml
    2 services:
    15 datanode:
    16   image: bde2020/hadoop-datanode:latest
    17   container_name: datanode
    18   environment:
    19     - CORE_CONF_fs_defaultFS=hdfs://namenode:8020
    20   volumes:
    21     - datanode-data:/hadoop/dfs/data
    22   depends_on:
    PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
    PS D:\xml\ParallelDistributed\Hadoop_11_25>
    History restored
    PS D:\xml\ParallelDistributed\Hadoop_11_25> docker-compose up -d
    time="2024-12-11T20:51:43+05:30" level=warning msg="D:\xml\ParallelDistributed\Hadoop_11_25\docker-compose.yml: 'version' is obsolete"
    [+] Running 3/3
    ✓ Container namenode Started 2.0s
    ✓ Container datanode Started 4.7s
    ✓ Container historyserver Started 7.7s
    PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -mkdir -p /input
    What's next:
    Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug namenode
    Learn more at https://docs.docker.com/go/debug-cli/
    PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -put /path/to/sample.txt /input
    put: '/path/to/sample.txt': No such file or directory
    What's next:
    Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug namenode
    Learn more at https://docs.docker.com/go/debug-cli/
```

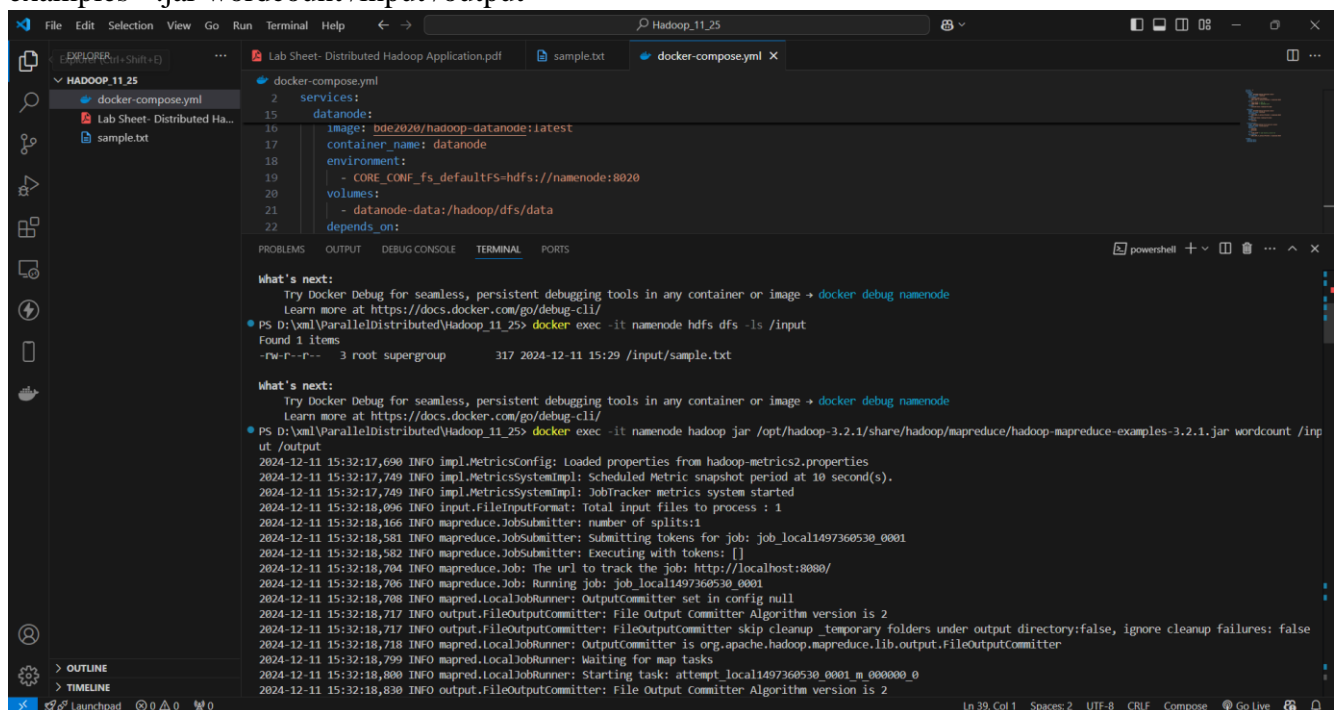
5. `docker exec -it namenode hdfs dfs -ls /input`



```
File Edit Selection View Go Run Terminal Help
Hadoop_11_25
EXPLORER
  HADOOP_11_25
    docker-compose.yml
    Lab Sheet- Distributed Ha...
    sample.txt
  docker-compose.yml
    2 services:
    15 datanode:
    16   image: bde2020/hadoop-datanode:latest
    17   container_name: datanode
    18   environment:
    19     - CORE_CONF_fs_defaultFS=hdfs://namenode:8020
    20   volumes:
    21     - datanode-data:/hadoop/dfs/data
    22   depends_on:
    PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
    PS D:\xml\ParallelDistributed\Hadoop_11_25>
    History restored
    PS D:\xml\ParallelDistributed\Hadoop_11_25> docker-compose up -d
    time="2024-12-11T20:51:43+05:30" level=warning msg="D:\xml\ParallelDistributed\Hadoop_11_25\docker-compose.yml: 'version' is obsolete"
    [+] Running 3/3
    ✓ Container namenode Started 2.0s
    ✓ Container datanode Started 4.7s
    ✓ Container historyserver Started 7.7s
    PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -mkdir -p /input
    What's next:
    Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug namenode
    Learn more at https://docs.docker.com/go/debug-cli/
    PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -put /path/to/sample.txt /input
    put: '/path/to/sample.txt': No such file or directory
    What's next:
    Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug namenode
    Learn more at https://docs.docker.com/go/debug-cli/
    PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -put ./sample.txt /input
    2024-12-11 15:29:01,993 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
    What's next:
    Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug namenode
    Learn more at https://docs.docker.com/go/debug-cli/
    PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -ls /input
    Found 1 items
    > OUTLINE
    > TIMELINE
    Launchpad 0 0 0 0
    Ln 39, Col 1 Spaces: 2 UTF-8 CRLF Compose Go Live
```



6. `docker exec -it namenode hadoop jar /opt/hadoop/share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar wordcount /input /output`




```
File Edit Selection View Go Run Terminal Help
Hadoop_11_25
EXPLORER
  HADOOP_11_25
    docker-compose.yml
    Lab Sheet- Distributed Ha...
    sample.txt
  docker-compose.yml
    15 services:
    16   datanode:
    17     image: bde2020/hadoop-datanode:latest
    18     container_name: datanode
    19     environment:
    20       - CORE_CONF_fs_defaultFS=hdfs://namenode:8020
    21     volumes:
    22       - datanode-data:/hadoop/dfs/data
    depends_on:

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
2024-12-11 15:32:19,211 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin: org.apache.hadoop.mapreduce.task.reduce.Shuffle@698e5fb6
2024-12-11 15:32:19,213 WARN impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2024-12-11 15:32:19,235 INFO reduce.MergeManagerImpl: MergerManager: memoryLimit=625370688, maxSingleShuffleLimit=156342672, mergeThreshold=412744672, ioSortFactor=10, memToMergeOutputsThreshold=10
2024-12-11 15:32:19,241 INFO reduce.EventFetcher: attempt_local1497360530_0001_r_000000_0 Thread started: EventFetcher for fetching Map Completion Events
2024-12-11 15:32:19,271 INFO reduce.LocalFetcher: localFetcher#1 about to shuffle output of map attempt_local1497360530_0001_r_000000_0 decomp: 547 len: 551 to MEMORY
2024-12-11 15:32:19,276 INFO reduce.InMemoryMapOutput: Read 547 bytes from map-output for attempt_local1497360530_0001_r_000000_0
2024-12-11 15:32:19,277 INFO reduce.MergeManagerImpl: closeInMemoryFile -> map-output of size: 547, inMemoryMapOutputs.size() -> 1, commitMemory -> 0, usedMemory -> 547
2024-12-11 15:32:19,279 INFO reduce.EventFetcher: EventFetcher is interrupted.. Returning
2024-12-11 15:32:19,282 INFO mapred.LocalJobRunner: 1 / 1 copied.
2024-12-11 15:32:19,282 INFO reduce.MergeManagerImpl: finalMerge called with 1 in-memory map-outputs and 0 on-disk map-outputs
2024-12-11 15:32:19,292 INFO mapred.Merger: Merging 1 sorted segments
2024-12-11 15:32:19,293 INFO reduce.MergeManagerImpl: Merged 1 segments, 547 bytes to disk to satisfy reduce memory limit
2024-12-11 15:32:19,295 INFO reduce.MergeManagerImpl: Merging 1 files, 551 bytes from disk
2024-12-11 15:32:19,298 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce
2024-12-11 15:32:19,298 INFO mapred.Merger: Merging 1 sorted segments
2024-12-11 15:32:19,299 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 539 bytes
2024-12-11 15:32:19,300 INFO mapred.LocalJobRunner: 1 / 1 copied.
2024-12-11 15:32:19,369 INFO Configuration.deprecation: mapred.skip.on is deprecated. Instead, use mapreduce.job.skiprecords
2024-12-11 15:32:19,416 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2024-12-11 15:32:19,480 INFO mapred.Task: Task:attempt_local1497360530_0001_r_000000_0 is done. And is in the process of committing
2024-12-11 15:32:19,484 INFO mapred.LocalJobRunner: 1 / 1 copied.
2024-12-11 15:32:19,484 INFO mapred.Task: Task attempt_local1497360530_0001_r_000000_0 is allowed to commit now
2024-12-11 15:32:19,514 INFO output.FileOutputCommitter: Saved output of task 'attempt_local1497360530_0001_r_000000_0' to hdfs://namenode:8020/output
2024-12-11 15:32:19,516 INFO mapred.LocalJobRunner: reduce > reduce
2024-12-11 15:32:19,516 INFO mapred.Task: Task 'attempt_local1497360530_0001_r_000000_0' done.
2024-12-11 15:32:19,516 INFO mapred.Task: Final Counters for attempt_local1497360530_0001_r_000000_0: Counters: 30
Ln 39, Col 1 Spaces: 2 UTF-8 CRLF Compose Go Live
```

```
File Edit Selection View Go Run Terminal Help
Hadoop_11_25
EXPLORER
  HADOOP_11_25
    docker-compose.yml
    Lab Sheet- Distributed Ha...
    sample.txt
  docker-compose.yml
    15 services:
    16   datanode:
    17     image: bde2020/hadoop-datanode:latest
    18     container_name: datanode
    19     environment:
    20       - CORE_CONF_fs_defaultFS=hdfs://namenode:8020
    21     volumes:
    22       - datanode-data:/hadoop/dfs/data
    depends_on:

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
2024-12-11 15:32:19,516 INFO mapred.Task: Final Counters for attempt_local1497360530_0001_r_000000_0: Counters: 30
File System Counters
  FILE: Number of bytes read=317829
  FILE: Number of bytes written=843053
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=317
  HDFS: Number of bytes written=377
  HDFS: Number of read operations=18
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=3
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Combine input records=0
  Combine output records=0
  Reduce input groups=42
  Reduce shuffle bytes=551
  Reduce input records=42
  Reduce output records=42
  Spilled Records=42
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=8
  Total committed heap usage (bytes)=267911168
Shuffle Errors
  BAD_ID=0
Ln 39, Col 1 Spaces: 2 UTF-8 CRLF Compose Go Live
```

```
File Edit Selection View Go Run Terminal Help
Hadoop_11_25
EXPLORER
  HADOOP_11_25
    docker-compose.yml
    Lab Sheet- Distributed Ha...
    sample.txt
  docker-compose.yml
    2 services:
    15   datanode:
    16     image: bde2020/hadoop-datanode:latest
    17     container_name: datanode
    18     environment:
    19       - CORE_CONF_fs_defaultsFS=hdfs://namenode:8020
    20     volumes:
    21       - datanode-data:/hadoop/dfs/data
    22     depends_on:

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
powershell + - - - - x

CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Output Format Counters
  Bytes Written=377
2024-12-11 15:32:19,517 INFO mapred.LocalJobRunner: Finishing task: attempt_local1497360530_0001_r_000000_0
2024-12-11 15:32:19,517 INFO mapred.LocalJobRunner: reduce task executor complete.
2024-12-11 15:32:19,732 INFO mapreduce.job: Job job_local1497360530_0001 running in uber mode : false
2024-12-11 15:32:19,733 INFO mapreduce.job: map 100% reduce 100%
2024-12-11 15:32:19,735 INFO mapreduce.job: Job job_local1497360530_0001 completed successfully
2024-12-11 15:32:19,745 INFO mapreduce.job: Counters: 36
File System Counters
  FILE: Number of bytes read=634524
  FILE: Number of bytes written=1685555
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=634
  HDFS: Number of bytes written=377
  HDFS: Number of read operations=15
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=4
  HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Map input records=3
  Map output records=47
```

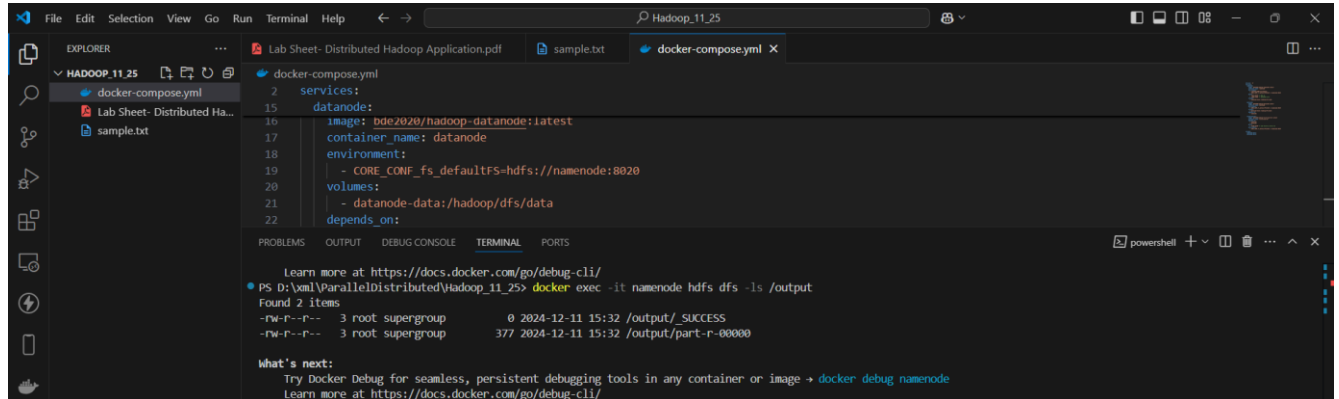
```
File Edit Selection View Go Run Terminal Help
Hadoop_11_25
EXPLORER
  HADOOP_11_25
    docker-compose.yml
    Lab Sheet- Distributed Ha...
    sample.txt
  docker-compose.yml
    2 services:
    15   datanode:
    16     image: bde2020/hadoop-datanode:latest
    17     container_name: datanode
    18     environment:
    19       - CORE_CONF_fs_defaultsFS=hdfs://namenode:8020
    20     volumes:
    21       - datanode-data:/hadoop/dfs/data
    22     depends_on:

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
powershell + - - - - x

Combine output records=42
Reduce input groups=42
Reduce shuffle bytes=551
Reduce input records=42
Reduce output records=42
Spilled Records=84
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=8
Total committed heap usage (bytes)=534773760
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=317
File Output Format Counters
  Bytes Written=377

What's next:
  Try Docker Debug for seamless, persistent debugging tools in any container or image → docker debug namenode
  Learn more at https://docs.docker.com/go/debug-cli/
• PS D:\vml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -ls /output
Found 2 items
```


7. docker exec -it namenode hdfs dfs -ls /output



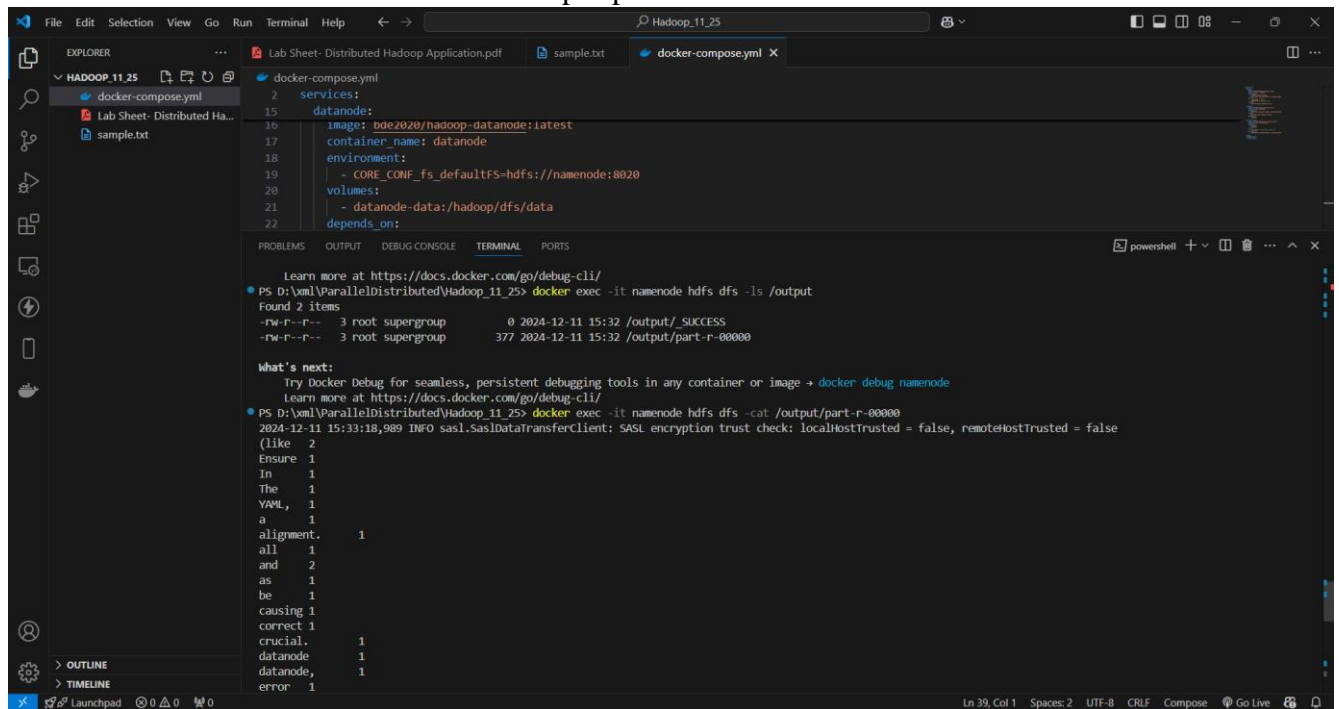
The screenshot shows a VS Code editor with a Docker Compose file for a Hadoop cluster. The file defines a service named 'datanode' using the 'bde2020/hadoop-datanode:latest' image. The environment is configured with 'CORE_CONF_fs_defaultFS=hdfs://namenode:8020' and a volume 'datanode-data' is mapped to '/hadoop/dfs/data'. The terminal output shows the command 'docker exec -it namenode hdfs dfs -ls /output' being executed successfully, listing two items: '0 2024-12-11 15:32 /output/_SUCCESS' and '377 2024-12-11 15:32 /output/part-r-000000'.

```
docker-compose.yml
2 services:
15   datanode:
16     image: bde2020/hadoop-datanode:latest
17     container_name: datanode
18     environment:
19       - CORE_CONF_fs_defaultFS=hdfs://namenode:8020
20     volumes:
21       - datanode-data:/hadoop/dfs/data
22     depends_on:

PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -ls /output
Found 2 items
-rw-r--r-- 3 root supergroup 0 2024-12-11 15:32 /output/_SUCCESS
-rw-r--r-- 3 root supergroup 377 2024-12-11 15:32 /output/part-r-000000

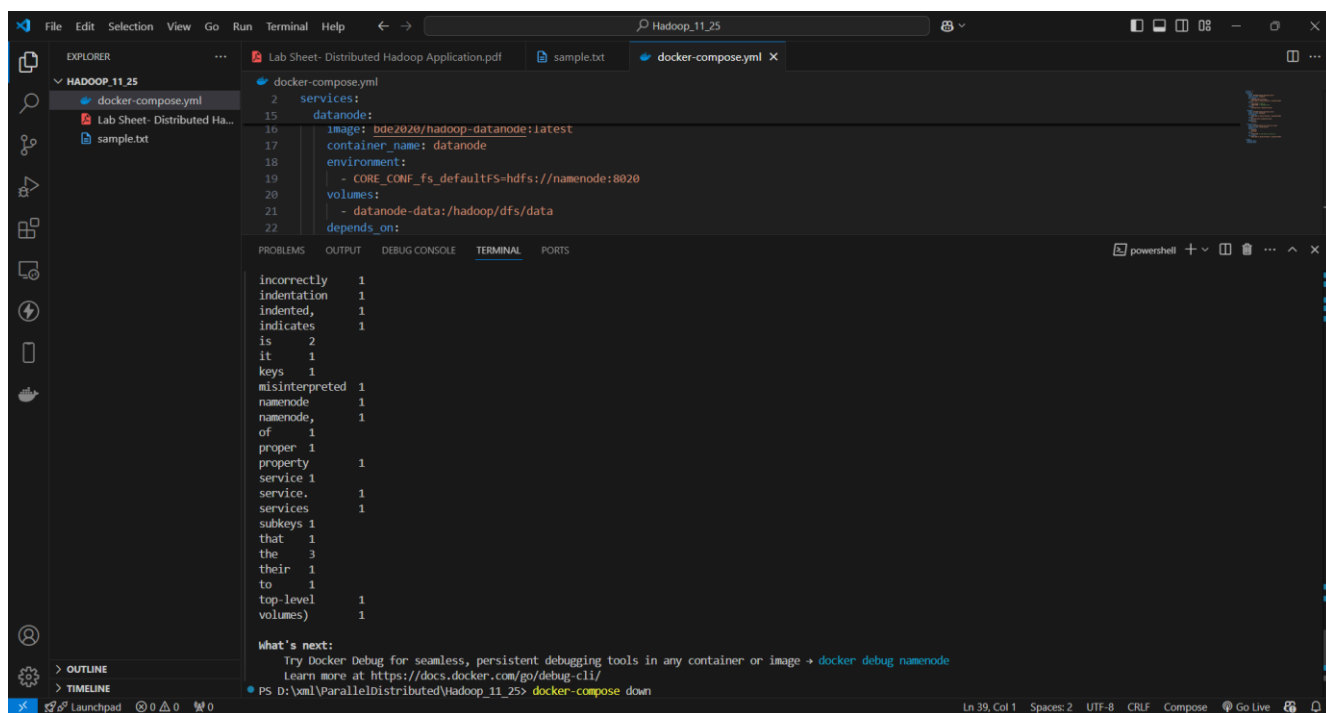
What's next:
Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug namenode
Learn more at https://docs.docker.com/go/debug-cli/
```

8. docker exec -it namenode hdfs dfs -cat /output/part-r-000000



The screenshot shows the same VS Code editor with the Docker Compose file. The terminal output shows the command 'docker exec -it namenode hdfs dfs -cat /output/part-r-000000' being executed successfully. The output of the cat command is displayed, showing a list of words and their counts, such as '2024-12-11 15:33:18,989 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false', '(like 2', 'Ensure 1', 'In 1', 'The 1', 'YAML, 1', 'a 1', 'alignment. 1', 'all 1', 'and 2', 'as 1', 'be 1', 'causing 1', 'correct 1', 'crucial. 1', 'datanode. 1', and 'error 1'.

```
PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -cat /output/part-r-000000
2024-12-11 15:33:18,989 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
(like 2
Ensure 1
In 1
The 1
YAML, 1
a 1
alignment. 1
all 1
and 2
as 1
be 1
causing 1
correct 1
crucial. 1
datanode. 1
error 1
```



9. docker-compose down

